

```

import os
import numpy as np
import pandas as pd
import random
import cv2
import matplotlib.pyplot as plt
%matplotlib inline
import keras as K
from keras.models import Model, Sequential
from keras.layers import Input, Dense, Flatten, Dropout, BatchNormalization
from keras.layers import Conv2D, SeparableConv2D, MaxPool2D, LeakyReLU, Activation
"from keras.optimizers import Adam"
from keras.preprocessing.image import ImageDataGenerator
from keras.callbacks import ModelCheckpoint, ReduceLROnPlateau, EarlyStopping
import tensorflow as tf
from mlxtend.plotting import plot_confusion_matrix

```

```

seed =232
np.random.seed(seed)
tf.random.set_seed(seed)


```

```

from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.



```
input_path = "/content/drive/MyDrive/AI-Project/chest_xray/"
```

```


fig, ax = plt.subplots(2,3, figsize=(15, 7))
ax = ax.ravel()
plt.tight_layout()

```

```

for i, _set in enumerate(['train', 'val', 'test']):
    set_path =input_path+_set
    ax[i].imshow(plt.imread(set_path+'/NORMAL/'+os.listdir(set_path+'/NORMAL')[0]),cmap='gray')
    ax[i].set_title('Set: {}, Condition: Normal'.format(_set))
    ax[i+3].imshow(plt.imread(set_path+'/PNEUMONIA/'+os.listdir(set_path+'/PNEUMONIA')[0]))
    ax[i+3].set_title('Set: {}, Condition: Pneumonia'.format(_set))

```



```
for _set in ['train', 'val', 'test']:
    n_normal = len(os.listdir(input_path + _set + '/NORMAL'))
    n_infect = len(os.listdir(input_path + _set + '/PNEUMONIA'))
    print('Set: {}, normal images: {}, pneumonia images: {}'.format(_set, n_normal, n_infe
```

```
Set: train, normal images: 1341, pneumonia images: 3875
```

```
Set: val, normal images: 8, pneumonia images: 8
```

```
Set: test, normal images: 234, pneumonia images: 390
```

```
def process_data(img_dims, batch_size):
```

```
# Data generation objects
```

```
train_datagen = ImageDataGenerator(rescale=1./255, zoom_range=0.3, vertical_flip=True)
```

```
test_val_datagen = ImageDataGenerator(rescale=1./255)
```

```
# This is fed to the network in the specified batch sizes and image dimensions
```

```
train_gen = train_datagen.flow_from_directory(
```

```
    directory=input_path+'train',
```

```
    target_size=(img_dims, img_dims),
```

```
    batch_size=batch_size,
```

```
    class_mode='binary',
```

```
    shuffle=True)
```

```
test_gen = test_val_datagen.flow_from_directory(
```

```
    directory=input_path+'test',
```

```
    target_size=(img_dims, img_dims),
```

```
    batch_size=batch_size,
```

```
    class_mode='binary',
```

```
    shuffle=True)
```

```
# I will be making predictions off of the test set in one batch size
```

```
# This is useful to be able to get the confusion matrix
```

```
test_data = []
```

```
test_labels = []
```

```
for cond in ['/NORMAL/', '/PNEUMONIA/']:
```

```
    for img in (os.listdir(input_path + 'test' + cond)):
```

```
        img = plt.imread(input_path+'test'+cond+img)
```

```

img = cv2.resize(img, (img_dims, img_dims))
img = np.dstack([img, img, img])
img = img.astype('float') / 255
if cond=='/NORMAL/':
    label = 0;
elif cond=='/PNEUMONIA/':
    label = 1;
test_data.append(img)
test_labels.append(label)

test_data = np.array(test_data)
test_labels = np.array(test_labels)

return train_gen, test_gen, test_data, test_labels

img_dims = 150
epochs = 10
batch_size = 32

train_gen, test_gen, test_data, test_labels = process_data(img_dims, batch_size)

    Found 5216 images belonging to 2 classes.
    Found 624 images belonging to 2 classes.

inputs = Input(shape=(img_dims, img_dims, 3))

#First conv block
x = Conv2D(filters=16, kernel_size=(3,3), activation='relu', padding='same')(inputs)
x = Conv2D(filters=16, kernel_size=(3,3), activation='relu', padding='same')(x)
x = MaxPool2D(pool_size=(2,2))(x)

#Second conv block
x = SeparableConv2D(filters=32, kernel_size=(3,3), activation='relu', padding='same')(x)
x = SeparableConv2D(filters=32, kernel_size=(3,3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = MaxPool2D(pool_size=(2,2))(x)

#Third conv block
x = SeparableConv2D(filters=64, kernel_size=(3,3), activation='relu', padding='same')(x)
x = SeparableConv2D(filters=64, kernel_size=(3,3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = MaxPool2D(pool_size=(2,2))(x)

#Fourth conv block
x = SeparableConv2D(filters=128, kernel_size=(3,3), activation='relu', padding='same')(x)
x = SeparableConv2D(filters=128, kernel_size=(3,3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = MaxPool2D(pool_size=(2,2))(x)
x = Dropout(rate=0.2)(x)

#Fifth conv block
x = SeparableConv2D(filters=256, kernel_size=(3,3), activation='relu', padding='same')(x)
x = SeparableConv2D(filters=256, kernel_size=(3,3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)

```

```

x = MaxPool2D(pool_size=(2,2))(x)
x = Dropout(rate=0.2)(x)

#FC layer
x = Flatten()(x)
x = Dense(units=512, activation='relu')(x)
x = Dropout(rate=0.7)(x)
x = Dense(units=128, activation='relu')(x)
x = Dropout(rate=0.5)(x)
x = Dense(units=64, activation='relu')(x)
x = Dropout(rate=0.3)(x)

#Output Layer
output = Dense(units = 1, activation='sigmoid')(x)

#Creating model and compiling
model = Model(inputs=inputs, outputs=output)
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

#Callbacks
checkpoint = ModelCheckpoint(filepath='best_weights.hdf5', save_best_only=True, save_weights_only=True)
lr_reduce = ReduceLROnPlateau(monitor='val_loss', factor=0.3, patience=2, verbose=2, mode='min')
early_stop = EarlyStopping(monitor='val_loss', min_delta=0.1, patience=1, mode='min')

print(model.summary())

```

Model: "model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 150, 150, 3)]	0
conv2d_2 (Conv2D)	(None, 150, 150, 16)	448
conv2d_3 (Conv2D)	(None, 150, 150, 16)	2320
max_pooling2d_5 (MaxPooling2D)	(None, 75, 75, 16)	0
separable_conv2d_8 (SeparableConv2D)	(None, 75, 75, 32)	688
separable_conv2d_9 (SeparableConv2D)	(None, 75, 75, 32)	1344
batch_normalization_4 (Batch Normalization)	(None, 75, 75, 32)	128
max_pooling2d_6 (MaxPooling2D)	(None, 37, 37, 32)	0
separable_conv2d_10 (SeparableConv2D)	(None, 37, 37, 64)	2400
separable_conv2d_11 (SeparableConv2D)	(None, 37, 37, 64)	4736

batch_normalization_5 (Batch Normalization)	(None, 37, 37, 64)	256
max_pooling2d_7 (MaxPooling2D)	(None, 18, 18, 64)	0
separable_conv2d_12 (SeparableConv2D)	(None, 18, 18, 128)	8896
separable_conv2d_13 (SeparableConv2D)	(None, 18, 18, 128)	17664
batch_normalization_6 (Batch Normalization)	(None, 18, 18, 128)	512
max_pooling2d_8 (MaxPooling2D)	(None, 9, 9, 128)	0
dropout_5 (Dropout)	(None, 9, 9, 128)	0
separable_conv2d_14 (SeparableConv2D)	(None, 9, 9, 256)	34176
separable_conv2d_15 (SeparableConv2D)	(None, 9, 9, 256)	68096
batch_normalization_7 (Batch Normalization)	(None, 9, 9, 256)	1024

```
hist = model.fit_generator(
    train_gen, steps_per_epoch = train_gen.samples // batch_size,
    epochs=epochs, validation_data=test_gen,
    validation_steps=test_gen.samples // batch_size, callbacks=[checkpoint, lr_redu
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:4: UserWarning: `Model` after removing the cwd from sys.path.

```
Epoch 1/10
163/163 [=====] - 2710s 17s/step - loss: 0.3792 - accuracy:
Epoch 2/10
163/163 [=====] - 318s 2s/step - loss: 0.2742 - accuracy: 0
Epoch 3/10
163/163 [=====] - 318s 2s/step - loss: 0.2707 - accuracy: 0
Epoch 4/10
163/163 [=====] - 312s 2s/step - loss: 0.2445 - accuracy: 0
Epoch 5/10
163/163 [=====] - ETA: 0s - loss: 0.2388 - accuracy: 0.9114
Epoch 00005: ReduceLROnPlateau reducing learning rate to 0.0003000000142492354.
163/163 [=====] - 317s 2s/step - loss: 0.2388 - accuracy: 0
Epoch 6/10
163/163 [=====] - 315s 2s/step - loss: 0.1848 - accuracy: 0
Epoch 7/10
163/163 [=====] - ETA: 0s - loss: 0.1722 - accuracy: 0.9369
Epoch 00007: ReduceLROnPlateau reducing learning rate to 9.0000000427477062e-05.
163/163 [=====] - 318s 2s/step - loss: 0.1722 - accuracy: 0
Epoch 8/10
163/163 [=====] - 319s 2s/step - loss: 0.1610 - accuracy: 0
Epoch 9/10
163/163 [=====] - ETA: 0s - loss: 0.1490 - accuracy: 0.9456
Epoch 00009: ReduceLROnPlateau reducing learning rate to 2.700000040931627e-05.
163/163 [=====] - 316s 2s/step - loss: 0.1490 - accuracy: 0
```


Epoch 10/10

163/163 [=====] - 314s 2s/step - loss: 0.1472 - accuracy: 0



```
fig, ax = plt.subplots(1,2, figsize=(10,3))
ax = ax.ravel()

for i, met in enumerate(['accuracy', 'loss']):
    ax[i].plot(hist.history[met])
    ax[i].plot(hist.history['val_' + met])
    ax[i].set_title('Model{}'.format(met))
    ax[i].set_xlabel('epochs')
    ax[i].set_ylabel(met)
    ax[i].legend(['train', 'val'])
```



```
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
preds = model.predict(test_data)
```

```
accuracy = accuracy_score(test_labels, np.round(preds))*100
cm = confusion_matrix(test_labels, np.round(preds))
tn, fp, fn, tp = cm.ravel()
```


```
print('-----CONFUSION MATRIX-----')
```

```
fig, ax = plot_confusion_matrix(conf_mat=cm, figsize=(10,8), hide_ticks=True, cmap=plt.cm.
plt.show())
```

```
print('\n-----TRAIN METRIC-----')
print('Train accuracy: {}'.format(np.round((hist.history['accuracy'][-1])*100,2)))
```

```
print('\n-----TEST METRIC-----')
precision = tp/(tp+fp)*100
recall = tp/(tp+fn)*100
print('Accuracy: {}'.format(accuracy))
print('Precision: {}'.format(precision))
print('Recall: {}'.format(recall))
print('F1-score: {}'.format(2*precision*recall/(precision+recall)))
```

-----CONFUSION MATRIX-----



-----TRAIN METRIC-----

Train accuracy: 95.0

-----TEST METRIC-----

Accuracy: 88.46153846153845%
Precision: 85.33333333333334%
Recall: 98.46153846153847%
F1-score: 91.42857142857143

✓ 7s completed at 5:23 PM

